

Control Systems Theory

Mathematical Foundations for the Triple Inverted Pendulum on a Cart

State-Space Modeling · *Lagrangian Dynamics*

PID Control · *LQR Optimal Control* · *MPC*

Extended Kalman Filter · *Unscented Kalman Filter*

MATLAB / Simulink Implementation Reference

Contents

1	System Description and Notation	3
1.1	Physical Setup	3
1.2	Coordinate Convention	3
2	Lagrangian Dynamics and Equations of Motion	3
2.1	Lagrangian Formulation	3
2.1.1	Kinetic Energy	3
2.1.2	Potential Energy	4
2.1.3	Compact Matrix Form	4
3	Linearisation and State-Space Representation	4
3.1	Nonlinear State-Space Form	4
3.2	First-Order Taylor Linearisation	4
3.3	Output Equation	4
3.4	Controllability and Observability	4
4	PID Control	5
4.1	Classical PID Structure	5
4.2	Application to the Triple Pendulum	5
4.3	Limitations	5
4.4	Tuning Methods	5
4.4.1	Ziegler–Nichols (Frequency Response)	5
4.4.2	Root Locus / Pole Placement	5
5	LQR Optimal State-Feedback Control	6
5.1	Problem Formulation	6
5.2	Algebraic Riccati Equation (ARE)	6
5.3	Weighting Matrix Selection	6
5.4	LQR vs. Pole Placement	6
5.5	Guaranteed Robustness Margins	6
6	Model Predictive Control (MPC)	7
6.1	Concept	7
6.2	Finite-Horizon Problem	7
6.3	Discrete-Time Model	7
6.4	QP Reformulation	7
6.5	MPC vs. LQR	7
7	State Estimation: Extended Kalman Filter	8
7.1	Motivation	8
7.2	EKF Algorithm	8
7.2.1	Prediction Step	8
7.2.2	Update Step	8
7.3	Limitations of the EKF	8
8	State Estimation: Unscented Kalman Filter	8
8.1	Unscented Transform	8
8.1.1	Sigma Points	9
8.1.2	Weights	9
8.2	UKF Algorithm	9
8.2.1	Prediction Step	9

8.2.2	Update Step	9
8.3	UKF vs. EKF: Summary	10
9	LQG: Combining LQR and the Kalman Filter	10
9.1	Separation Principle	10
9.2	Closed-Loop Architecture	10
10	Stability Analysis	10
10.1	Lyapunov Stability	10
10.2	Open-Loop Instability	10
11	MATLAB Implementation Reference	11
11.1	Script Execution Order	11
11.2	Key MATLAB Commands	11
12	Glossary	11

1 System Description and Notation

1.1 Physical Setup

The system consists of a cart of mass M free to slide along a horizontal rail, with three rigid pendulum links attached in series. Link i has mass m_i and length ℓ_i , for $i = 1, 2, 3$. The cart is actuated by a horizontal force u .

State Vector (interleaved convention):

$$\mathbf{x} = \begin{bmatrix} x \\ \dot{x} \\ \theta_1 \\ \dot{\theta}_1 \\ \theta_2 \\ \dot{\theta}_2 \\ \theta_3 \\ \dot{\theta}_3 \end{bmatrix} \in \mathbb{R}^8$$

where x is cart displacement and θ_i is the absolute angle of link i from vertical.

1.2 Coordinate Convention

All angles θ_i are measured from the upright vertical. The equilibrium of interest is:

$$\mathbf{x}^* = \mathbf{0}_{8 \times 1}, \quad u^* = 0.$$

This equilibrium is *open-loop unstable*: any infinitesimal perturbation causes the links to fall.

2 Lagrangian Dynamics and Equations of Motion

2.1 Lagrangian Formulation

The equations of motion are derived via the Euler–Lagrange equations:

$$\frac{d}{dt} \frac{\partial \mathcal{L}}{\partial \dot{q}_k} - \frac{\partial \mathcal{L}}{\partial q_k} = Q_k,$$

where $\mathcal{L} = T - V$ is the Lagrangian (kinetic minus potential energy), $\mathbf{q} = [x, \theta_1, \theta_2, \theta_3]^\top$ are the generalised coordinates, and Q_k are generalised forces.

2.1.1 Kinetic Energy

The position of the centre of mass of link i (length ℓ_i , assumed uniform) is:

$$x_{c,i} = x + \sum_{j=1}^{i-1} \ell_j \sin \theta_j + \frac{\ell_i}{2} \sin \theta_i, \quad y_{c,i} = \sum_{j=1}^{i-1} \ell_j \cos \theta_j + \frac{\ell_i}{2} \cos \theta_i.$$

The total kinetic energy is:

$$T = \frac{1}{2} M \dot{x}^2 + \sum_{i=1}^3 \frac{1}{2} m_i (\dot{x}_{c,i}^2 + \dot{y}_{c,i}^2) + \sum_{i=1}^3 \frac{1}{2} I_i \dot{\theta}_i^2,$$

where $I_i = \frac{1}{12} m_i \ell_i^2$ is the moment of inertia about the link's centre of mass.

2.1.2 Potential Energy

$$V = \sum_{i=1}^3 m_i g y_{c,i}.$$

2.1.3 Compact Matrix Form

The equations of motion take the standard manipulator form:

$$\mathbf{M}(\mathbf{q}) \ddot{\mathbf{q}} + \mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \dot{\mathbf{q}} + \mathbf{G}(\mathbf{q}) = \mathbf{B} u,$$

where:

- $\mathbf{M}(\mathbf{q}) \in \mathbb{R}^{4 \times 4}$ is the symmetric positive-definite inertia matrix,
- $\mathbf{C}(\mathbf{q}, \dot{\mathbf{q}}) \in \mathbb{R}^{4 \times 4}$ contains Coriolis and centrifugal terms,
- $\mathbf{G}(\mathbf{q}) \in \mathbb{R}^4$ is the gravity vector,
- $\mathbf{B} = [1, 0, 0, 0]^\top$ maps the scalar force u to generalised coordinates.

3 Linearisation and State-Space Representation

3.1 Nonlinear State-Space Form

Defining $\mathbf{x} \in \mathbb{R}^8$ as above and $\mathbf{u} = u \in \mathbb{R}$, the dynamics are written as:

$$\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x}, u), \quad \mathbf{y} = \mathbf{h}(\mathbf{x}),$$

where $\mathbf{f} : \mathbb{R}^8 \times \mathbb{R} \rightarrow \mathbb{R}^8$ encodes the full nonlinear ODEs.

3.2 First-Order Taylor Linearisation

Expanding $\mathbf{f}$ around the equilibrium $(\mathbf{x}^*, u^*) = (0, 0)$:

$$\dot{\mathbf{x}} \approx \mathbf{A}\mathbf{x} + \mathbf{B}u,$$

$$\mathbf{A} = \left. \frac{\partial \mathbf{f}}{\partial \mathbf{x}} \right|_{\mathbf{x}^*, u^*} \in \mathbb{R}^{8 \times 8}, \quad \mathbf{B} = \left. \frac{\partial \mathbf{f}}{\partial u} \right|_{\mathbf{x}^*, u^*} \in \mathbb{R}^{8 \times 1}.$$

At the upright equilibrium, $\cos \theta_i \approx 1$ and $\sin \theta_i \approx \theta_i$, so the inertia matrix reduces to a constant and all Coriolis terms vanish.

3.3 Output Equation

If only the cart position and three angles are measured:

$$\mathbf{y} = \mathbf{C}\mathbf{x}, \quad \mathbf{C} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{bmatrix} \in \mathbb{R}^{4 \times 8}.$$

3.4 Controllability and Observability

Definition 3.1. The system $(\mathbf{A}, \mathbf{B})$ is *controllable* if and only if the controllability matrix

$$\mathcal{C} = [\mathbf{B} \quad \mathbf{A}\mathbf{B} \quad \mathbf{A}^2\mathbf{B} \quad \cdots \quad \mathbf{A}^7\mathbf{B}] \in \mathbb{R}^{8 \times 8}$$

has rank 8.

Definition 3.2. The system (C, A) is *observable* if and only if the observability matrix

$$\mathcal{O} = [C^\top \quad A^\top C^\top \quad \dots \quad (A^\top)^7 C^\top]^\top \in \mathbb{R}^{32 \times 8}$$

has rank 8.

Both conditions *must* hold before any model-based controller or observer can be designed. In MATLAB: `rank(ctrb(A,B))` and `rank(observ(A,C))`.

4 PID Control

4.1 Classical PID Structure

A PID controller generates a control signal from the error $e(t) = r(t) - y(t)$:

$$u(t) = K_p e(t) + K_i \int_0^t e(\tau) d\tau + K_d \frac{de(t)}{dt},$$

with transfer function:

$$C(s) = K_p + \frac{K_i}{s} + K_d s = \frac{K_d s^2 + K_p s + K_i}{s}.$$

4.2 Application to the Triple Pendulum

For the triple pendulum, a PID loop is applied per-link using the angle error $e_i = -\theta_i$:

$$u_i(t) = K_{p,i} \theta_i + K_{d,i} \dot{\theta}_i + K_{i,i} \int_0^t \theta_i d\tau.$$

The total force command is the superposition:

$$u(t) = \sum_{i=1}^3 u_i(t).$$

4.3 Limitations

Why PID is insufficient here:

1. Ignores coupling between links: each loop treats other states as disturbances.
2. Cart position x is not naturally regulated by angle-only feedback.
3. Tuning 9+ parameters for 3 interacting loops is non-systematic.
4. No formal optimality guarantee.

LQR addresses all four limitations through full-state feedback.

4.4 Tuning Methods

4.4.1 Ziegler–Nichols (Frequency Response)

1. Set $K_i = K_d = 0$. Increase K_p until sustained oscillation occurs at gain K_u and period T_u .
2. Apply: $K_p = 0.6K_u$, $K_i = 1.2K_u/T_u$, $K_d = 0.075K_u T_u$.

4.4.2 Root Locus / Pole Placement

Design $C(s)$ so that closed-loop poles lie in a desired region of the left half-plane. For a second-order target with damping ζ and natural frequency ω_n , desired poles are $s = -\zeta\omega_n \pm j\omega_n\sqrt{1-\zeta^2}$.

5 LQR Optimal State-Feedback Control

5.1 Problem Formulation

Given the linear system $\dot{\mathbf{x}} = \mathbf{A}\mathbf{x} + \mathbf{B}u$, find the state-feedback gain $\mathbf{K} \in \mathbb{R}^{1 \times 8}$ such that $u = -\mathbf{K}\mathbf{x}$ minimises the infinite-horizon quadratic cost:

$$J = \int_0^\infty (\mathbf{x}^\top \mathbf{Q} \mathbf{x} + u^\top \mathbf{R} u) dt,$$

where:

- $\mathbf{Q} = \mathbf{Q}^\top \succeq 0 \in \mathbb{R}^{8 \times 8}$ penalises state deviation,
- $\mathbf{R} = \mathbf{R}^\top \succ 0 \in \mathbb{R}^{1 \times 1}$ penalises control effort.

5.2 Algebraic Riccati Equation (ARE)

Theorem 5.1 (LQR Optimality). If $(\mathbf{A}, \mathbf{B})$ is stabilisable and $(\mathbf{A}, \mathbf{Q}^{1/2})$ is detectable, there exists a unique symmetric positive-semidefinite solution $\mathbf{P}$ to the ARE:

$$\mathbf{A}^\top \mathbf{P} + \mathbf{P} \mathbf{A} - \mathbf{P} \mathbf{B} \mathbf{R}^{-1} \mathbf{B}^\top \mathbf{P} + \mathbf{Q} = \mathbf{0}.$$

The optimal gain is then:

$$\mathbf{K}^* = \mathbf{R}^{-1} \mathbf{B}^\top \mathbf{P}.$$

The closed-loop system $\dot{\mathbf{x}} = (\mathbf{A} - \mathbf{B}\mathbf{K}^*)\mathbf{x}$ is asymptotically stable.

5.3 Weighting Matrix Selection

A standard starting point is *Bryson's rule*:

$$Q_{ii} = \frac{1}{x_{i,\max}^2}, \quad R = \frac{1}{u_{\max}^2},$$

where $x_{i,\max}$ and $u_{\max}$ are the maximum acceptable values for state i and the control input.

Tuning intuition:

- Increase Q_{ii} to penalise excursions in state i more heavily.
- Increase R to reduce control effort (slower response, less actuator wear).
- The ratio $\mathbf{Q}/R$ governs the speed-effort tradeoff.

5.4 LQR vs. Pole Placement

For an 8-state system, placing all 8 closed-loop poles manually is impractical: small changes in desired pole locations can yield wildly different gain values. LQR provides a *systematic, physically interpretable alternative* via the $\mathbf{Q}$ and $\mathbf{R}$ matrices.

5.5 Guaranteed Robustness Margins

For a SISO LQR design, the optimal gain has guaranteed phase margin $\geq 60^\circ$ and gain margin $[\frac{1}{2}, \infty)$. This is a fundamental property of LQR that PID does not inherit automatically.

6 Model Predictive Control (MPC)

6.1 Concept

MPC solves a finite-horizon optimal control problem *at every time step*, applies only the first control action, then re-optimises at the next step (receding horizon). This allows:

- Explicit handling of state and input constraints.
- Preview of future reference trajectories.
- Online adaptation to model changes.

6.2 Finite-Horizon Problem

At time t , given current state $\hat{\mathbf{x}}_t$, solve:

$$\min_{\{u_k\}_{k=0}^{N-1}} \sum_{k=0}^{N-1} \left(\mathbf{x}_k^\top \mathbf{Q} \mathbf{x}_k + u_k^\top \mathbf{R} u_k \right) + \mathbf{x}_N^\top \mathbf{P} \mathbf{x}_N$$

subject to:

$$\mathbf{x}_{k+1} = \mathbf{A}_d \mathbf{x}_k + \mathbf{B}_d u_k, \quad \mathbf{x}_k \in \mathcal{X}, \quad u_k \in \mathcal{U},$$

where $\mathbf{A}_d, \mathbf{B}_d$ are the discrete-time system matrices (via zero-order hold with sample time T_s), N is the prediction horizon, $\mathbf{P}$ is the LQR terminal cost, and $\mathcal{X}, \mathcal{U}$ are polyhedral constraint sets.

6.3 Discrete-Time Model

Discretise the continuous system with sample period T_s using the matrix exponential:

$$\mathbf{A}_d = e^{\mathbf{A}T_s}, \quad \mathbf{B}_d = \int_0^{T_s} e^{\mathbf{A}\tau} \mathbf{B} d\tau \approx (\mathbf{A}^{-1} (e^{\mathbf{A}T_s} - \mathbf{I})) \mathbf{B}.$$

6.4 QP Reformulation

Stacking predictions over the horizon yields the condensed QP:

$$\min_{\mathbf{U}} \frac{1}{2} \mathbf{U}^\top \mathbf{H} \mathbf{U} + \mathbf{f}^\top \mathbf{U} \quad \text{s.t.} \quad \mathbf{G} \mathbf{U} \leq \mathbf{w} + \mathbf{S} \hat{\mathbf{x}}_t,$$

where $\mathbf{U} = [u_0, \dots, u_{N-1}]^\top \in \mathbb{R}^N$, and $\mathbf{H}, \mathbf{f}, \mathbf{G}, \mathbf{S}$ are derived from $\mathbf{Q}, \mathbf{R}, \mathbf{P}, \mathbf{A}_d, \mathbf{B}_d$. This QP is solved at every sample with a dedicated solver (e.g. OSQP).

6.5 MPC vs. LQR

Property	LQR	MPC
State constraints	No	Yes
Input constraints	No	Yes
Computational cost	Low (offline)	High (online QP)
Horizon	Infinite	Finite N
Reference preview	No	Yes
Optimality	Global (unconstrained)	Local (constrained)

7 State Estimation: Extended Kalman Filter

7.1 Motivation

In practice, only a subset of states is measurable and measurements are corrupted by noise:

$$\dot{\mathbf{x}} = f(\mathbf{x}, u) + \mathbf{L}\mathbf{w}, \quad \mathbf{y} = h(\mathbf{x}) + \mathbf{v},$$

where $\mathbf{w} \sim \mathcal{N}(\mathbf{0}, \mathbf{Q}_w)$ is process noise and $\mathbf{v} \sim \mathcal{N}(\mathbf{0}, \mathbf{R}_v)$ is measurement noise.

7.2 EKF Algorithm

The EKF handles nonlinear f and h by linearising at each step.

7.2.1 Prediction Step

$$\hat{\mathbf{x}}_{k|k-1} = f(\hat{\mathbf{x}}_{k-1|k-1}, u_{k-1}), \quad (1)$$

$$\mathbf{P}_{k|k-1} = \mathbf{F}_{k-1} \mathbf{P}_{k-1|k-1} \mathbf{F}_{k-1}^\top + \mathbf{L} \mathbf{Q}_w \mathbf{L}^\top, \quad (2)$$

where $\mathbf{F}_{k-1} = \left. \frac{\partial f}{\partial \mathbf{x}} \right|_{\hat{\mathbf{x}}_{k-1|k-1}}$ is the Jacobian.

7.2.2 Update Step

$$\mathbf{S}_k = \mathbf{H}_k \mathbf{P}_{k|k-1} \mathbf{H}_k^\top + \mathbf{R}_v, \quad (3)$$

$$\mathbf{K}_k = \mathbf{P}_{k|k-1} \mathbf{H}_k^\top \mathbf{S}_k^{-1}, \quad (4)$$

$$\hat{\mathbf{x}}_{k|k} = \hat{\mathbf{x}}_{k|k-1} + \mathbf{K}_k (\mathbf{y}_k - h(\hat{\mathbf{x}}_{k|k-1})), \quad (5)$$

$$\mathbf{P}_{k|k} = (\mathbf{I} - \mathbf{K}_k \mathbf{H}_k) \mathbf{P}_{k|k-1}, \quad (6)$$

where $\mathbf{H}_k = \left. \frac{\partial h}{\partial \mathbf{x}} \right|_{\hat{\mathbf{x}}_{k|k-1}}$.

7.3 Limitations of the EKF

- Jacobian linearisation discards higher-order terms; accuracy degrades for highly nonlinear systems.
- May diverge if initial estimate or covariance is far from the truth.
- Requires analytical Jacobians (or numerical differentiation).
- The propagated covariance can become non-positive-definite in poorly conditioned cases.

8 State Estimation: Unscented Kalman Filter

8.1 Unscented Transform

The UKF replaces Jacobian linearisation with the *unscented transform*: a set of $2n+1$ deterministically chosen *sigma points* that capture the mean and covariance of a Gaussian distribution exactly to the *third order* of a Taylor expansion. For $n = 8$ states:

8.1.1 Sigma Points

Given mean $\hat{\mathbf{x}}$ and covariance $\mathbf{P}$, compute the matrix square root $\mathbf{S} = \text{chol}((n + \lambda)\mathbf{P})$:

$$\chi_0 = \hat{\mathbf{x}},$$

$$\chi_i = \hat{\mathbf{x}} + \mathbf{S}_i, \quad i = 1, \dots, n,$$

$$\chi_{n+i} = \hat{\mathbf{x}} - \mathbf{S}_i, \quad i = 1, \dots, n,$$

where $\mathbf{S}_i$ denotes the i -th column of $\mathbf{S}$ and $\lambda = \alpha^2(n + \kappa) - n$.

8.1.2 Weights

$$W_0^{(m)} = \frac{\lambda}{n + \lambda}, \quad W_0^{(c)} = \frac{\lambda}{n + \lambda} + (1 - \alpha^2 + \beta),$$

$$W_i^{(m)} = W_i^{(c)} = \frac{1}{2(n + \lambda)}, \quad i = 1, \dots, 2n.$$

Typical values: $\alpha = 10^{-3}$, $\beta = 2$, $\kappa = 0$.

8.2 UKF Algorithm

8.2.1 Prediction Step

1. Generate $2n + 1$ sigma points from $\hat{\mathbf{x}}_{k-1|k-1}$, $\mathbf{P}_{k-1|k-1}$.
2. Propagate each sigma point through the nonlinear dynamics: $\chi_i^- = f(\chi_i, u_{k-1})$.
3. Reconstruct predicted mean and covariance:

$$\hat{\mathbf{x}}_{k|k-1} = \sum_{i=0}^{2n} W_i^{(m)} \chi_i^-,$$

$$\mathbf{P}_{k|k-1} = \sum_{i=0}^{2n} W_i^{(c)} (\chi_i^- - \hat{\mathbf{x}}_{k|k-1})(\chi_i^- - \hat{\mathbf{x}}_{k|k-1})^\top + \mathbf{Q}_w.$$

8.2.2 Update Step

1. Generate new sigma points from $\hat{\mathbf{x}}_{k|k-1}$, $\mathbf{P}_{k|k-1}$.
2. Propagate through measurement function: $\gamma_i = h(\chi_i^-)$.
3. Compute predicted measurement mean: $\hat{\mathbf{y}}_k = \sum_{i=0}^{2n} W_i^{(m)} \gamma_i$.
4. Compute innovation covariance and cross-covariance:

$$\mathbf{S}_k = \sum_{i=0}^{2n} W_i^{(c)} (\gamma_i - \hat{\mathbf{y}}_k)(\gamma_i - \hat{\mathbf{y}}_k)^\top + \mathbf{R}_v,$$

$$\mathbf{T}_k = \sum_{i=0}^{2n} W_i^{(c)} (\chi_i^- - \hat{\mathbf{x}}_{k|k-1})(\gamma_i - \hat{\mathbf{y}}_k)^\top.$$

5. Kalman gain and update:

$$\mathbf{K}_k = \mathbf{T}_k \mathbf{S}_k^{-1}, \quad \hat{\mathbf{x}}_{k|k} = \hat{\mathbf{x}}_{k|k-1} + \mathbf{K}_k (\mathbf{y}_k - \hat{\mathbf{y}}_k), \quad \mathbf{P}_{k|k} = \mathbf{P}_{k|k-1} - \mathbf{K}_k \mathbf{S}_k \mathbf{K}_k^\top.$$

8.3 UKF vs. EKF: Summary

Property	EKF	UKF
Nonlinearity handling	First-order linearisation	Third-order (sigma points)
Jacobians required	Yes	No
Accuracy	Lower for strong nonlinearities	Higher
Computational cost	$\mathcal{O}(n^3)$ (Cholesky)	$\mathcal{O}(n^3)$ (comparable)
Implementation complexity	Medium	Medium–High
Divergence risk	Higher	Lower

Why UKF for the triple pendulum: The nonlinear dynamics involve products of $\sin \theta_i$, $\cos \theta_i$, and $\dot{\theta}_i^2$. For large angles or fast transients (especially open-loop), the first-order approximation in the EKF introduces significant error. The UKF’s sigma-point approach propagates the distribution more accurately without requiring Jacobians.

9 LQG: Combining LQR and the Kalman Filter


9.1 Separation Principle

When the system is linear and noise is Gaussian, the optimal control and estimation problems *decouple*:

- Design the LQR gain $\mathbf{K}$ ignoring noise.
- Design the Kalman filter gain $\mathbf{L}$ ignoring control.
- Combine: $u = -\mathbf{K}\hat{\mathbf{x}}$.

The resulting LQG controller is optimal. For nonlinear systems (triple pendulum), the separation principle is applied approximately by replacing the Kalman filter with an EKF or UKF.

9.2 Closed-Loop Architecture



10 Stability Analysis

10.1 Lyapunov Stability

Theorem 10.1 (Lyapunov’s Second Method). The equilibrium $\mathbf{x} = \mathbf{0}$ of $\dot{\mathbf{x}} = \mathbf{f}(\mathbf{x})$ is asymptotically stable if there exists $V : \mathbb{R}^n \rightarrow \mathbb{R}$ with:

$$V(\mathbf{0}) = 0, \quad V(\mathbf{x}) > 0 \quad \forall \mathbf{x} \neq \mathbf{0}, \quad \dot{V}(\mathbf{x}) = \nabla V \cdot \mathbf{f}(\mathbf{x}) < 0 \quad \forall \mathbf{x} \neq \mathbf{0}.$$

For the LQR-controlled linear system, the quadratic function $V(\mathbf{x}) = \mathbf{x}^\top \mathbf{P} \mathbf{x}$ serves as a Lyapunov function, with $\dot{V} = -\mathbf{x}^\top \mathbf{Q}^* \mathbf{x} < 0$, confirming global asymptotic stability of the closed-loop linear system.

10.2 Open-Loop Instability

The open-loop eigenvalues of $\mathbf{A}$ include right-half-plane (RHP) values — the system has at least one positive real eigenvalue per unstable link. LQR moves all closed-loop eigenvalues of $\mathbf{A} - \mathbf{B}\mathbf{K}$ to the left half-plane.

11 MATLAB Implementation Reference

11.1 Script Execution Order

Step	Script	Function
1	parameters.m	Define physical constants
2	linearize.m	Compute $\mathbf{A}, \mathbf{B}, \mathbf{C}$; check rank
3	lqr_design.m	Solve ARE, compute $\mathbf{K}$; saves lqr_data.mat
4	open_loop_sim.m	Simulate unstable open-loop
5	pid_design.m	Manual PID tuning + simulation
6	ukf_design.m	UKF state estimation with noisy measurements
7	compare_controllers.m	LQR vs PID performance metrics
8	animate_pendulum.m	Generate GIF animations

11.2 Key MATLAB Commands

Command	Purpose
<code>[A,B] = linearize(params)</code>	Numerical Jacobian at upright eq.
<code>rank(ctrb(A,B))</code>	Verify controllability (must be 8)
<code>[K,P,e] = lqr(A,B,Q,R)</code>	Solve ARE, return gain and eigenvalues
<code>eig(A - B*K)</code>	Verify closed-loop stability
<code>ode45(@(t,x)odefun,tspan,x0)</code>	Nonlinear closed-loop simulation
<code>expm(A*Ts)</code>	Discrete-time $\mathbf{A}_d$ for MPC

12 Glossary

ARE	Algebraic Riccati Equation — governs the optimal LQR cost matrix $\mathbf{P}$.
Controllability	Ability to drive any state to any other state in finite time.
EKF	Extended Kalman Filter — nonlinear estimator using first-order Jacobian linearisation.
LQG	Linear Quadratic Gaussian — optimal output-feedback controller (LQR + Kalman filter).
LQR	Linear Quadratic Regulator — state-feedback controller minimising a quadratic cost.
MPC	Model Predictive Control — receding-horizon constrained optimisation.
Observability	Ability to reconstruct all states from output measurements alone.
PID	Proportional-Integral-Derivative controller — classical single-loop regulator.
UKF	Unscented Kalman Filter — nonlinear estimator using deterministic sigma points.
$\mathbf{Q}, \mathbf{R}$	LQR weighting matrices — trading state regulation against control effort.
$\mathbf{Q}_w, \mathbf{R}_v$	Kalman filter noise covariances — process and measurement noise.